

# Sentiment Analysis using Streaming Spark

BD\_078\_460\_474\_565

Anushka Hebbar  
Shree Thavarakere  
Vibha Masti  
Shruvi D

# Dataset

- Machine learning and spark streaming dataset
- It consists of 2 features:
  - Feature 0 - Sentiment (an integer - 0 or 4, which represents the label of the data)
  - Feature 1 - Tweet (the actual twitter message)
- The dataset is divided into train and test sets

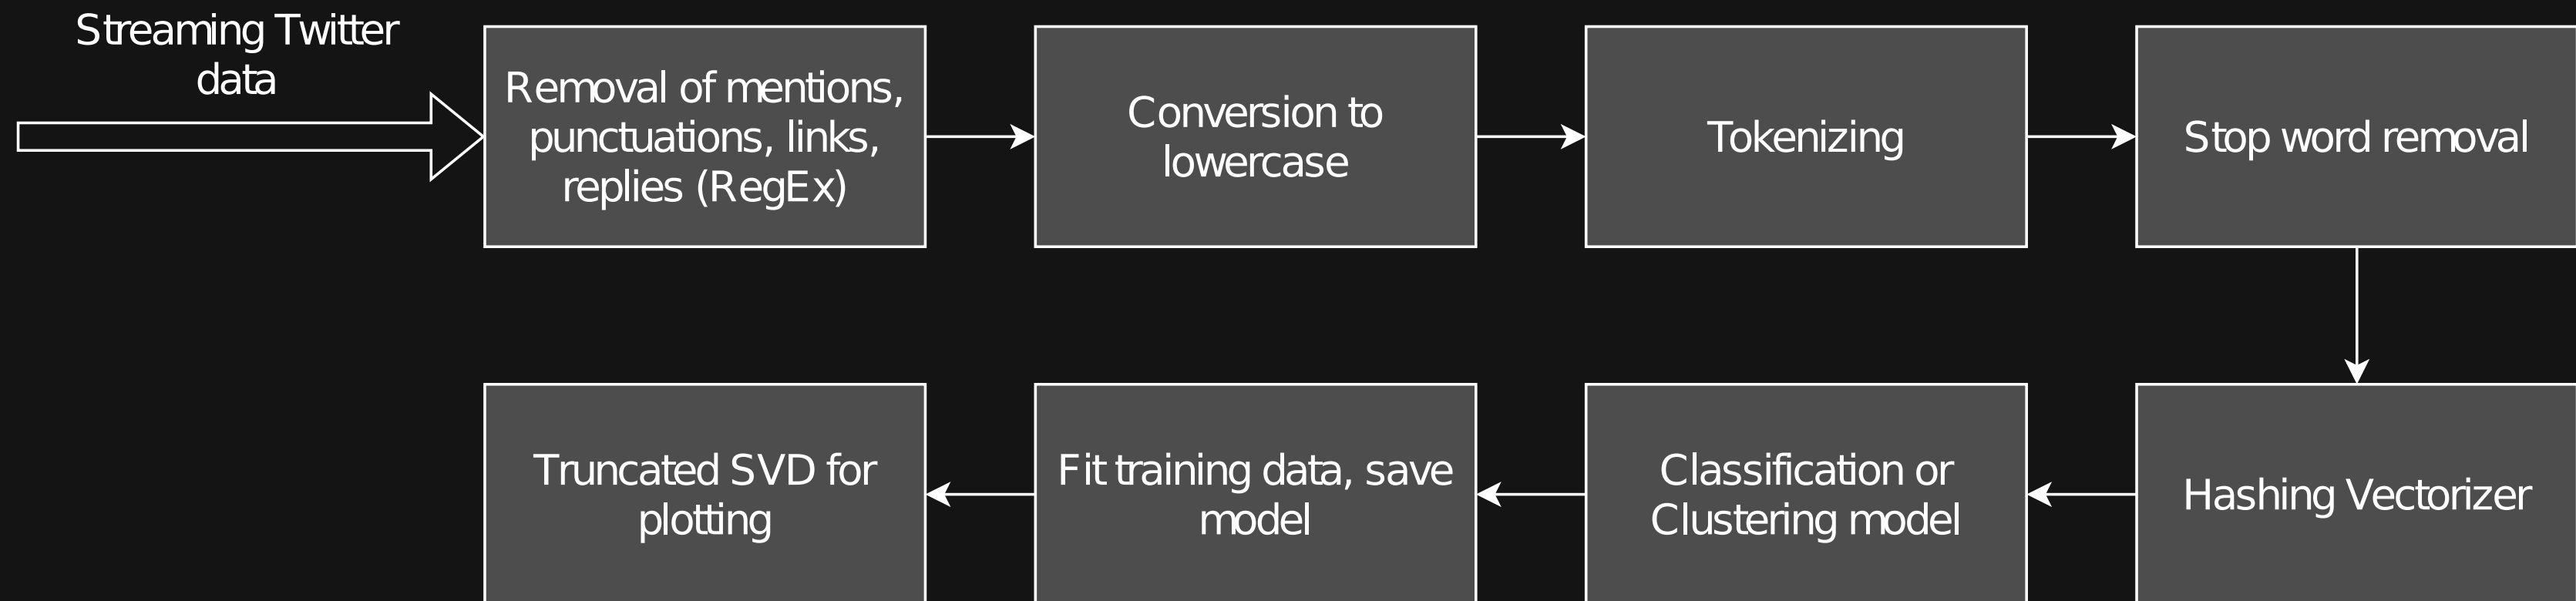
# Structure of Repository

- train.py - acts as main file for training functionality
- preprocessing
  - preprocess.py
- classification\_models
  - logistic\_regression.py
  - multinomial\_nb.py
  - passive\_aggressive.py
- clustering\_models
  - kmeans\_clustering.py
  - birch\_clustering.py
- test.py - acts as main file for testing functionality

# Preprocessing the Stream

- Data is read in as a dstream
- For each RDD in the stream a function called 'process' is called
- In process, first a dictionary is created from the result of `rdd.collect()`
- This dictionary is then converted to a spark dataframe, which is used for further steps in preprocessing

# NLP Pipeline



# Preprocessing the Data

## BASIC PREPROCESSING

- Removal of mentions, punctuations, links and replies are removed with Regular Expressions
- Conversion to lowercase

## TOKENIZER

- Each record converted to an array of individual words

## STOP WORDS REMOVER

- Stop words like 'a', 'the', 'on' etc removed

## HASHING VECTORIZER

- Conversion of collection of text documents to Sparse Matrix (L2 norm)
- Size of matrix remains constant and vocabulary not stored in memory
- 2-gram

## TRUNCATED SVD

- This is used for dimensionality reduction.

# Libraries Used

Apart from the built in libraries of python, the following libraries were used:

- pyspark
  - For all data storage, manipulation and processing
  - For streaming support
- sklearn
  - For all the classification and clustering models
  - For some preprocessing and vectorizing techniques.
  - For decomposition - dimensionality reduction techniques.
  - For metrics.
- numpy
  - For processing of arrays, vectors and sparse arrays (sparse arrays were passed to the models)
- matplotlib
  - For plotting clusters
  - For visualizing results of classification models

# Classification Models

## Stochastic Gradient Descent Classifiers:

This estimator implements regularized linear models with stochastic gradient descent (SGD) learning. SGD allows incremental learning via the `partial_fit` method.

The model it fits can be controlled with the `loss` parameter

In our implementation, we used the following models:

- Logistic Regression model - 'log'
- Support Vector model - 'hinge'
- Perceptron algorithm model - 'perceptron'



# Classification Models


## Multinomial Naive Bayes:

MultinomialNB implements the naive Bayes algorithm for multinomially distributed data, and is one of the two classic naive Bayes variants used in text classification.

The hyperparameters here are:

- alpha: Additive (Laplace/Lidstone) smoothing parameter (0 for no smoothing). In our implementation we tried values of
  - 1.0
  - 0.5
  - 0.7
- fit\_prior: Whether to learn class prior probabilities or not
  - We set this to True
- class\_prior: Prior probabilities of the classes
  - We set this to None (we didn't have any prior probabilities)

# Classification Models



## Passive Aggressive Classifier:

The passive-aggressive algorithms are a family of algorithms for large-scale learning. It is an 'online-learning' algorithm, which makes it perfect for training streaming data.

- Passive: If the prediction is correct, keep the model and do not make any changes. i.e., the data in the example is not enough to cause any changes in the model.
- Aggressive: If the prediction is incorrect, make changes to the model. i.e., some change to the model may correct it.
- Hyperparameter-  $C$  : This is the regularization parameter, and denotes the penalization the model will make on an incorrect prediction.
- In our implementation, we used the following values of  $C$ :
  - 0.2
  - 0.5
  - 1.0

# Clustering Models

## KMeans:

The KMeans algorithm clusters data by trying to separate samples in  $n$  groups of equal variance, minimizing a criterion known as the inertia or within-cluster sum-of-squares.

Parameters used are:

- *n\_clusters=2*
- *init='k-means++'* (selects initial cluster centers for k-mean clustering in a smart way to speed up convergence)
- *n\_init=2* (Number of time the k-means algorithm will be run with different centroid seeds)
  - Varying this from 1-2 gave a very slight improvement in accuracy, but further increase had no impact
- *init\_size=1000* (Number of samples to randomly sample for speeding up the initialization)
  - Varying this value from 0 -100-1000 did not have any impact on accuracy
- *verbose=False*
- *max\_iter=1000*
  - Varying this from 1000-5000 had no impact on accuracy

# Clustering Models

## **Birch Clustering - Tried, but not included in final project**

The Birch builds a tree called the Clustering Feature Tree (CFT) for the given data. The data is essentially lossy compressed to a set of Clustering Feature nodes (CF Nodes). The CF Nodes have a number of subclusters called Clustering Feature subclusters (CF Subclusters) .

The CF Subclusters hold the necessary information for clustering.

Parameters:

- $n\_clusters=2$  (Number of clusters after the final clustering step, which treats the subclusters from the leaves as new samples)

However, the efficiency of this algorithm was very less, and it took a very long time. Also, it did not give any significant improvement over KMeans. Therefore, it was not used in the final training.

# Results

Each model (3 variations of the 3 classification models and 1 clustering model) was tested on different batch sizes

**The batch sizes trained on:**

1. 2000

2. 2500

3. 3000


4. 4000

5. 5000

Accuracy:


$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

ROC curve: TPR vs FPR




# Clustering Plots


Batch size 2000



Batch size 2500




Batch size 3000






# Clustering Plots

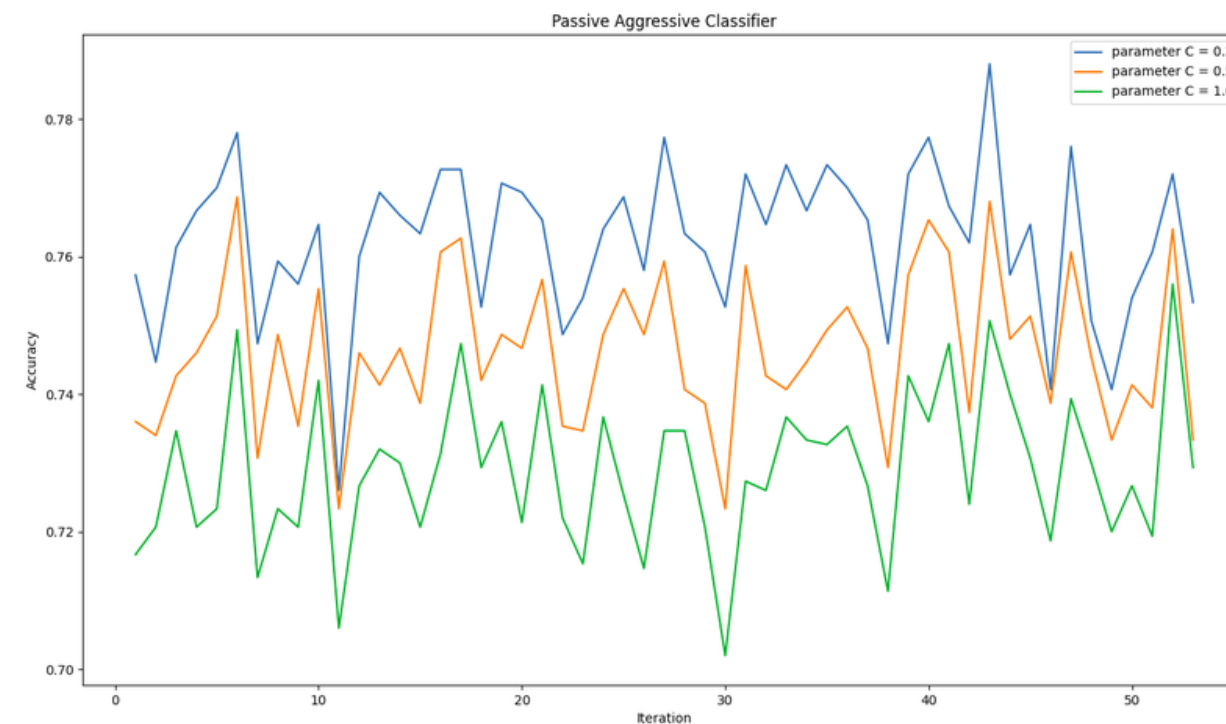
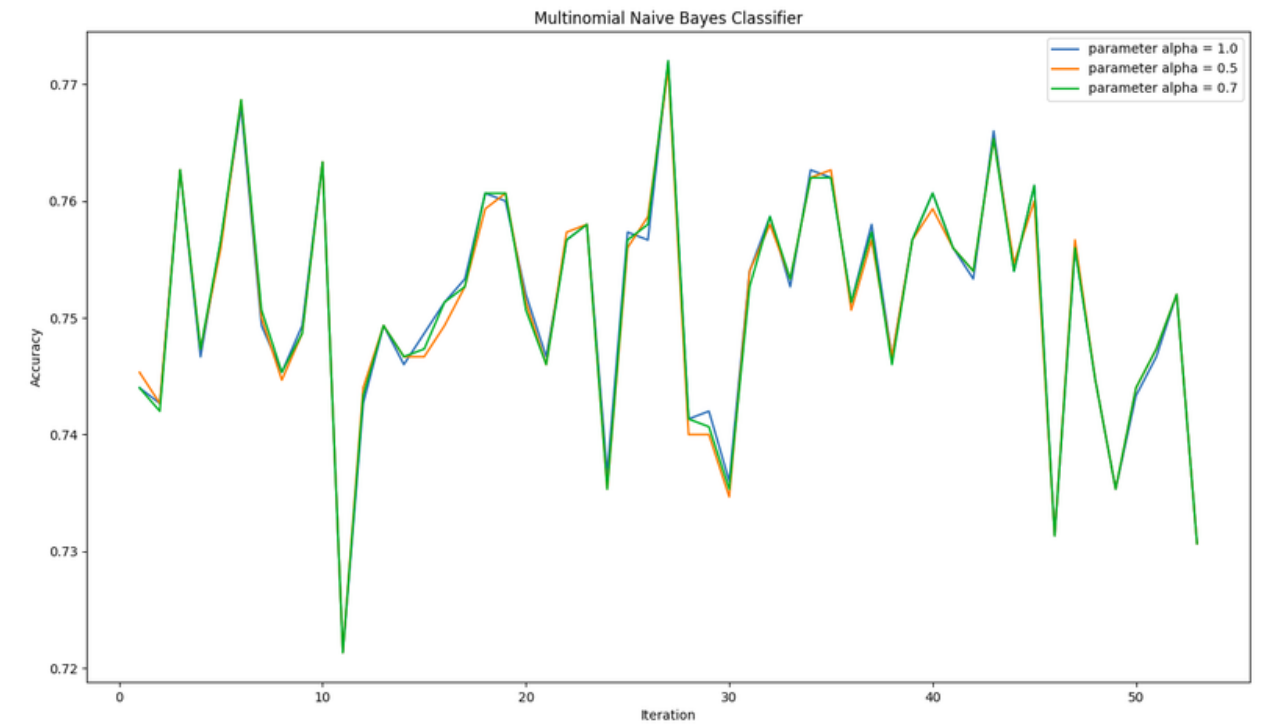

Batch size 4000



Batch size 5000





# Results - Hyperparameters




Batch Size 2000






# Accuracy - Hyperparameters




Batch Size 5000



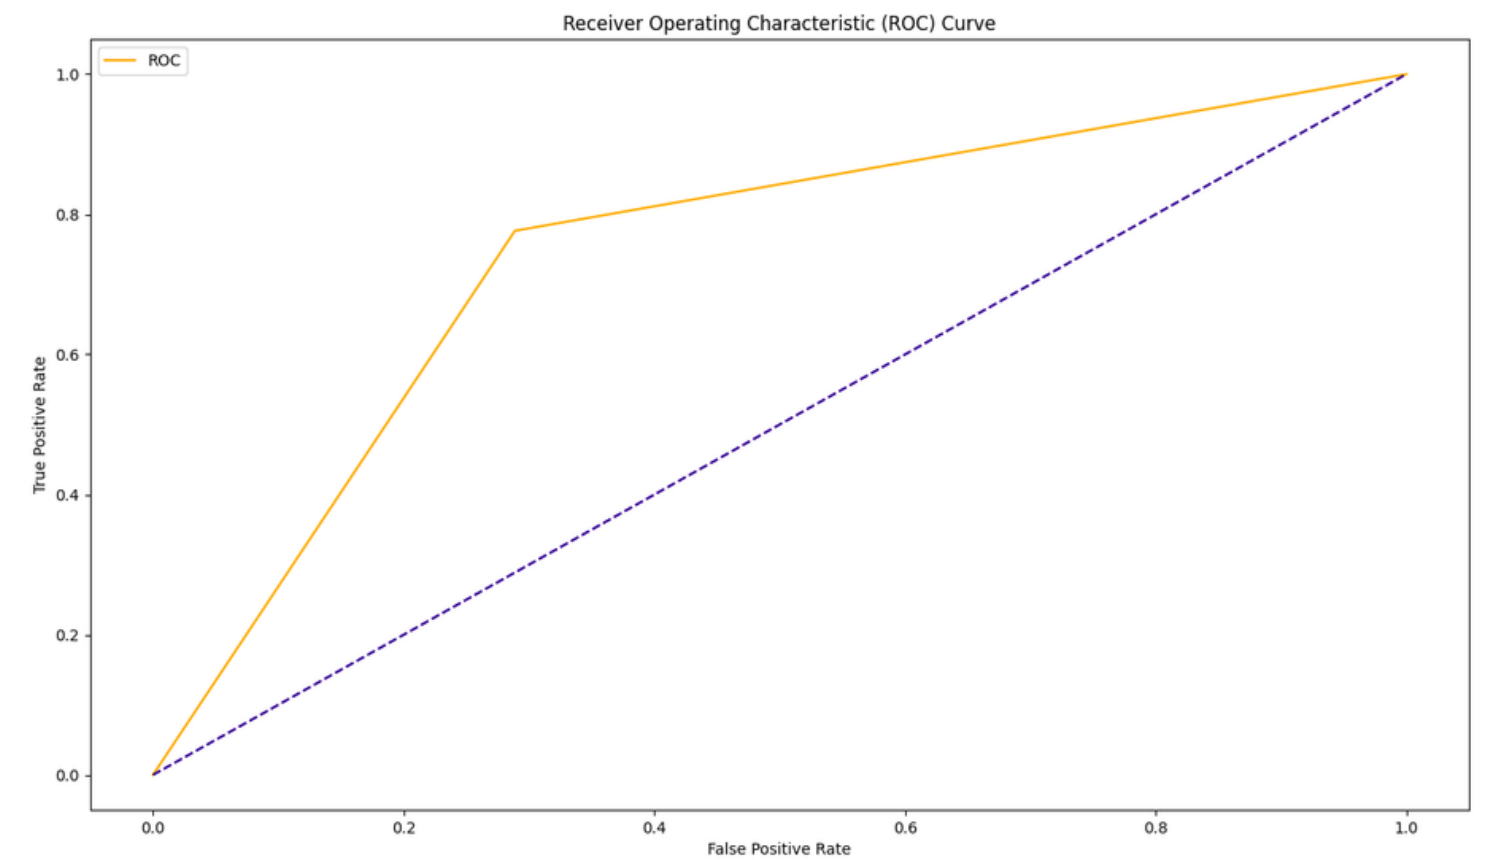
# Accuracy (Batch Sizes) - 3 models



# Results - ROC Curves (SGD)




Batch Size 2000 - First SGD Model

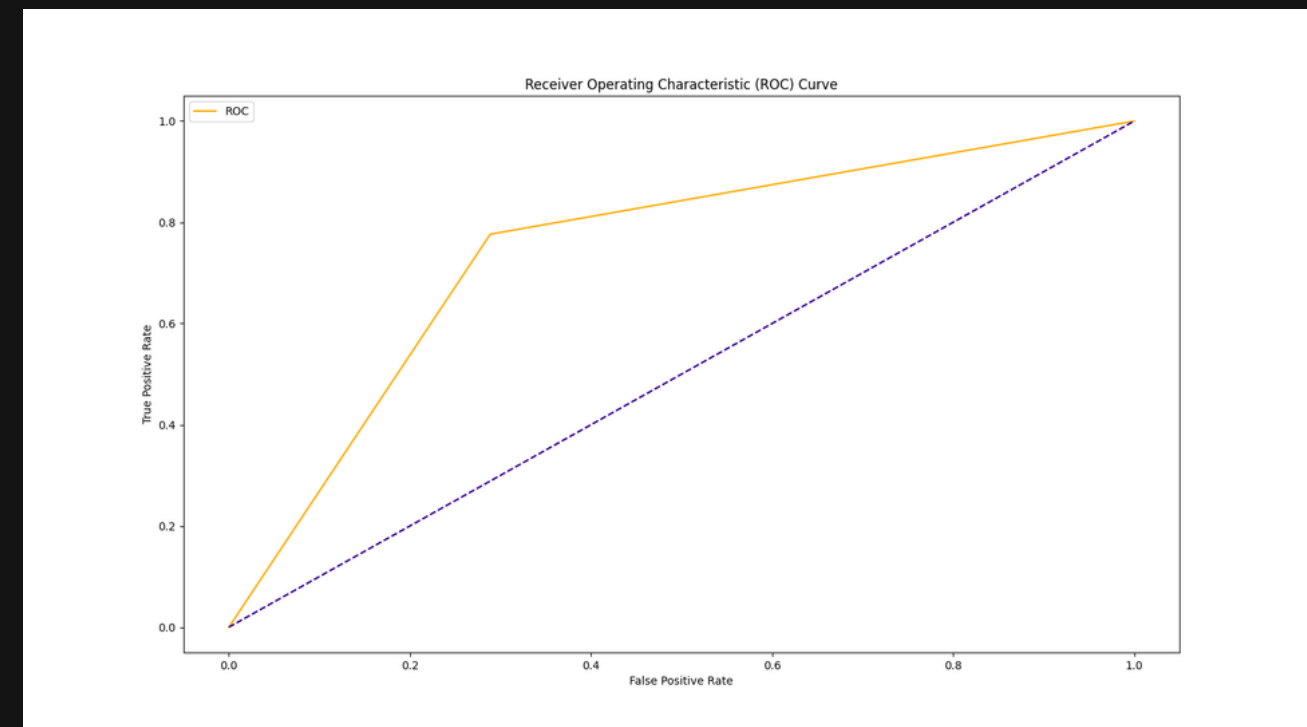


Batch Size 2500 - First SGD Model


# Results - ROC Curves (SGD)



Batch Size 3000 - First SGD Model



Batch Size 4000 - First SGD Model



Batch Size 5000 - First SGD Model

# Summarised Results

Model	Batch Size	Hyperparameter	Average Test Accuracy
SGD	2000	Log loss	0.7432
		<b>Hinge loss</b>	<b>0.7455</b>
		Perceptron loss	0.6872
	3000	Log loss	0.7431
		<b>Hinge loss</b>	<b>0.7453</b>
		Perceptron loss	0.6867
	4000	Log loss	0.7432
		<b>Hinge loss</b>	<b>0.7456</b>
		Perceptron loss	0.6983
	5000	Log loss	0.7432
		Hinge loss	0.7456
		Perceptron loss	0.6909

# Summarised Results

Model	Batch Size	Hyperparameter	Average Test Accuracy
MNB	2000	1.0	0.7511
		0.5	0.7509
		0.7	0.7501
	3000	1.0	0.7511
		0.5	0.7508
		0.7	0.7509
	4000	1.0	0.7511
		0.5	0.7509
		0.7	0.7511
	5000	1.0	0.7511
		0.5	0.7509
		0.7	0.7511

# Summarised Results

Model	Batch Size	Hyperparameter	Average Test Accuracy
PAC	2000	<b>0.2</b>	<b>0.7624</b>
		0.5	0.7463
		1.0	0.7289
	3000	<b>0.2</b>	<b>0.7623</b>
		0.5	0.7461
		1.0	0.7231
	4000	<b>0.2</b>	<b>0.7624</b>
		0.5	0.7461
		1.0	0.7247
	5000	<b>0.2</b>	<b>0.7614</b>
		0.5	0.7470
		1.0	0.7246

# Conclusion

- Accuracy did not change much when batch size varied from 2000-5000
  - Could not train smaller batch sizes on our systems
- LR and SVM performed similarly, better than perceptron
- Changing alpha in MNB had almost no impact
- Changing C in PAC had slight impact
  - Optimal C: 0.2
- All models - ~75% accuracy
- K-means clustering not effective
- Possible future improvements
  - DBScan clustering
  - Smaller batch sizes
  -



**Thank you**